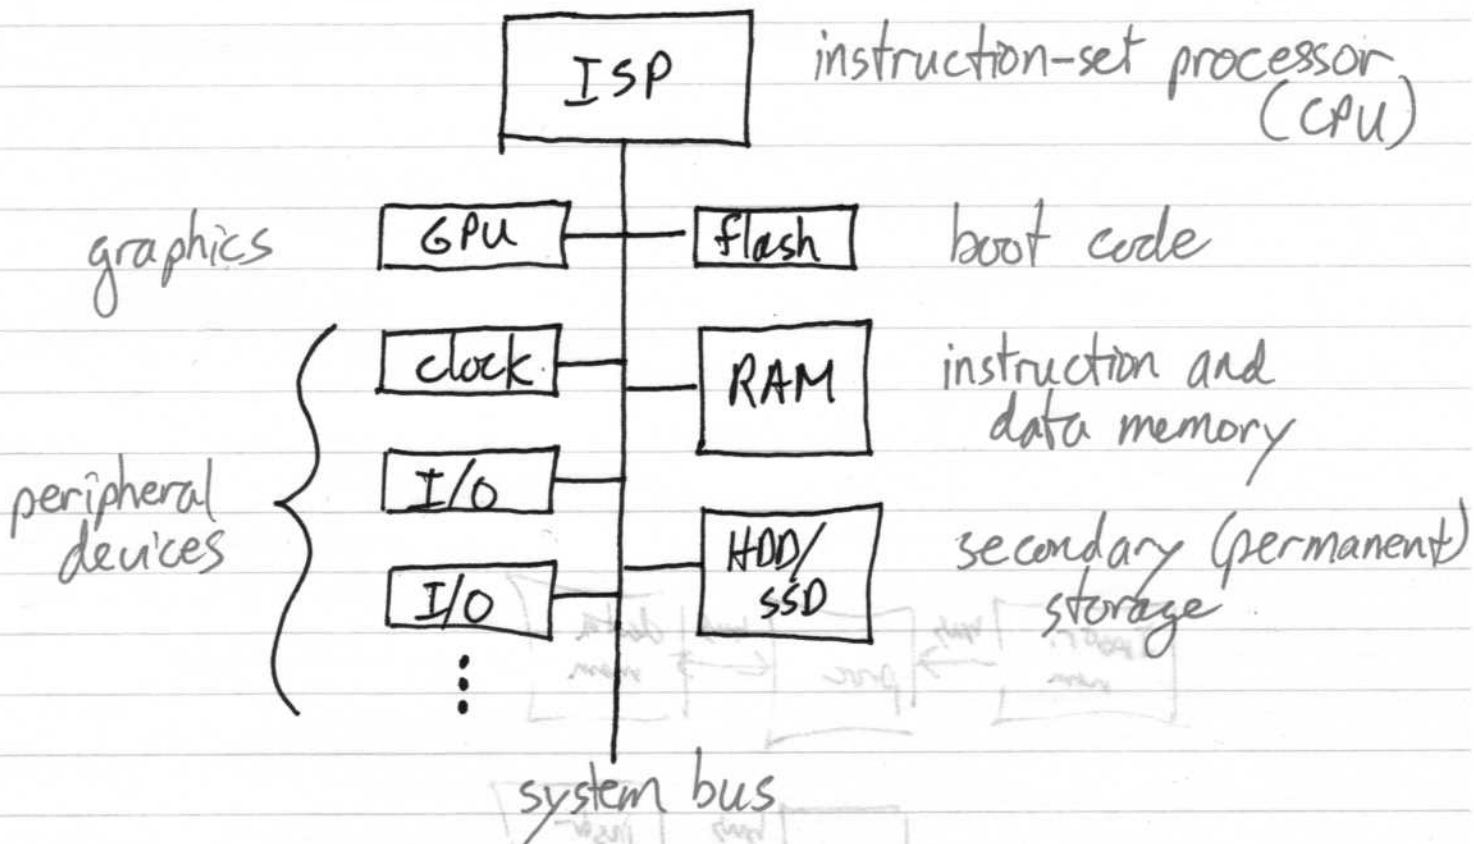


Computer Systems

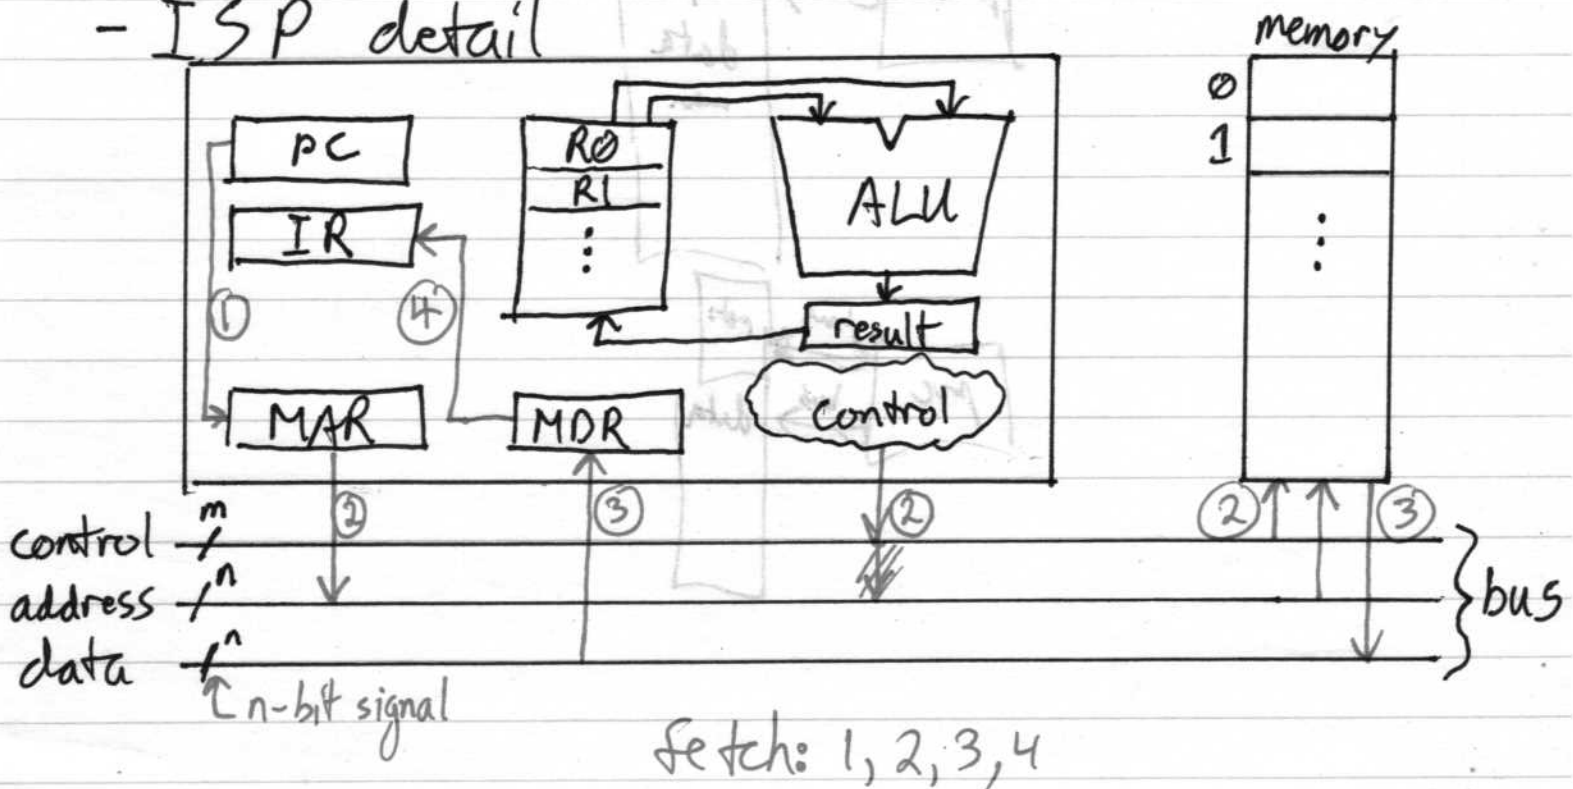
①

1) Structure

[3.2, 3.4]



- ISP detail



(2)

PC = Program Counter

- stores next instruction address

IR = Instruction Register

- holds current instruction for decoding

MAR = Memory Address Register

- outputs address to memory

MDR = Memory Data Register

- holds data to/from memory

Harvard Architecture

- separate memories for instructions and data

vs

Von Neumann Architecture

- one memory for instructions and data

- instruction execution

e.g. ADD R2, R1, #50

destination register
source operand register
immediate operand

F: Fetch (instruction)
and increment PC

MAR \leftarrow PC, bus read, IR \leftarrow MDR
PC \leftarrow PC + 4

D: decode (figure out
what to do)

IR	ADD	2	1	50
	opcode	dst reg #	src reg #	immed. data

C: compute

result \leftarrow R1 + 50

WB: writeback (result
to a register)

R2 \leftarrow result


(3)

- ISA : Instruction Set Architecture
 - describes processor functionality (instruction encoding, memory model, register set, exception handling)
 - abstraction layer between hw and sw

- ISA types

CISC


- complex instruction-set computer
- popular 1960s - mid 1980s
- variable-length instructions
- each instruction can do a lot e.g. string copy
- shorter programs
- multiple ~~inst~~ cycles per instruction



e.g. x86

RISC

- reduced instr.-set comp.
- introduced mid-80s
- fixed-length instructions
- simple instructions
- longer programs
- pipelined execution \Rightarrow 1 cycle per instr.



e.g. ARM 3-stage
IBM Power9 12-stages

2.1) Storage (Memory) Technology


- non-volatile memory retains data when powered off

e.g. ROM read-only memory

- contents are programmed at the foundry
- used for startup code (BIOS)

e.g. Flash reprogrammable memory

- can only be reprogrammed in blocks of 16-128 kiB
- also used for startup code (e.g. in LPC1768)




- volatile memory loses data when powered off

e.g. SRAM static RAM

- a fast, expensive memory
- used in caches

uses an inverter pair



e.g. DRAM dynamic RAM

- a slower, cheaper memory
- used for main memory

uses a capacitor




6

- hard disk drives (HDD)

- data is stored on spinning magnetic platters

- the read/write head seeks to the correct track, then waits for the right sector to spin under it



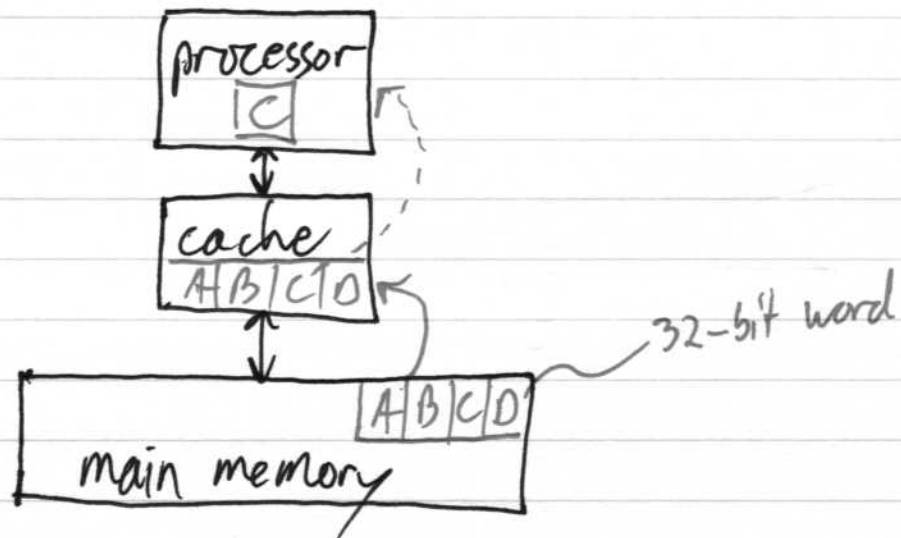
- data is stored in sectors of 512 B or 4 KiB

- HDDs are very slow (~ 10 ms latency) but have large capacity and high bandwidth

2.2) Cache Memories

- main memory is large but has high latency (~ 100 s clock cycles)
- caches are smaller, faster memories storing a subset of data from main memory

e.g.




- based on the idea that if a processor needs data from memory:
 - a) it will need the same data again soon
(temporal locality)
 - b) it will need other data nearby
(spatial locality)
- operation:
 - stores recently used data
 - retrieves data from memory in blocks (this takes advantage of memory's high bandwidth)

8

- cache blocks are "tagged" with their memory address
- for every processor memory access:
 - it will check the tags for a matching address
 - if found "hit" - returns data
 - if not found "miss" - retrieve block from memory, then return the data
- there are often multiple levels of cache
 - e.g. L1, L2, L3
 - smaller caches are faster (L1)
 - larger caches have higher hit rates (L3)

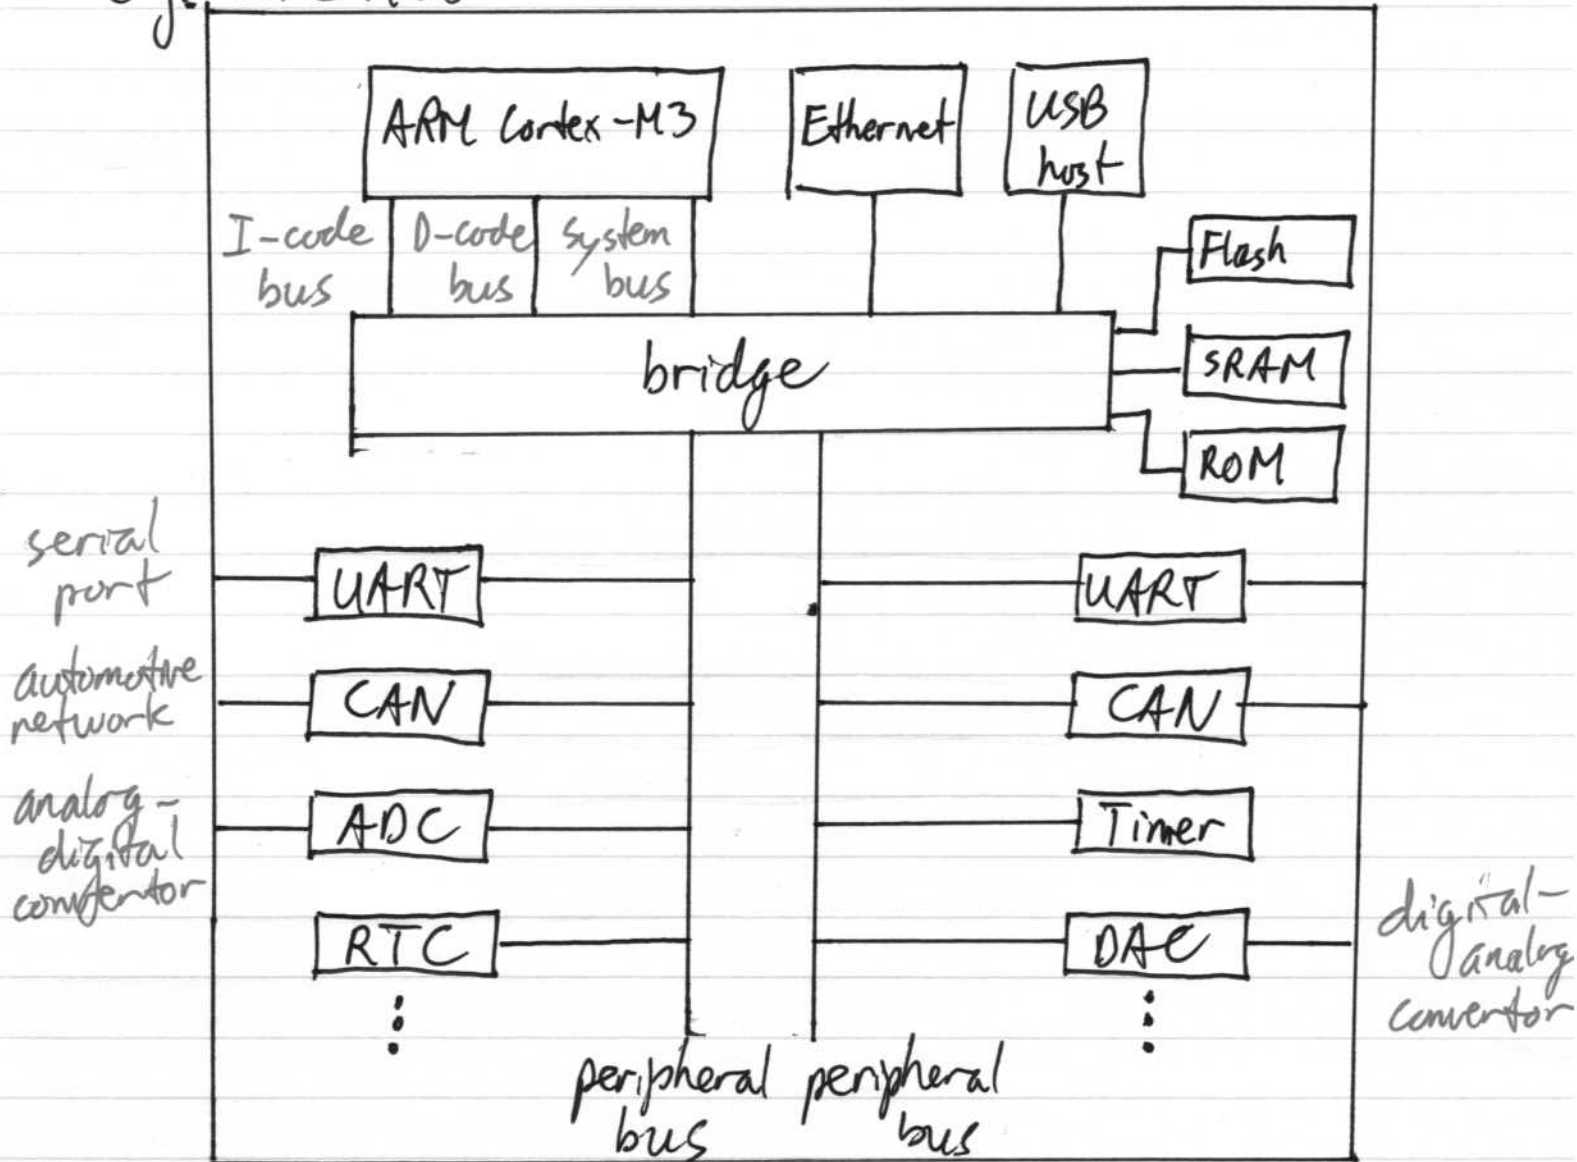
e.g. dual-core



3) Microcontroller Unit (MCU)

- an MCU integrates the processor, memory, real-time clock (RTC), and I/O devices on one chip

e.g. LPC1768



- I-code bus: for instructions (connects to Flash)
 - D-code bus: for data (connects to SRAM)
 - System bus: for I/O
- } Harvard Architecture

- ARM Cortex-M3 processor

- word size = 32 bits (all registers are 32 bits)
- it can address 4 GiB of memory ($2^{32} = 4 \text{ GiB}$)
- it has 16 general-purpose registers (R0-R15);
R15 is the PC
- it can do 8-bit, 16-bit and 32-bit data transfers
- no cache