

I/O

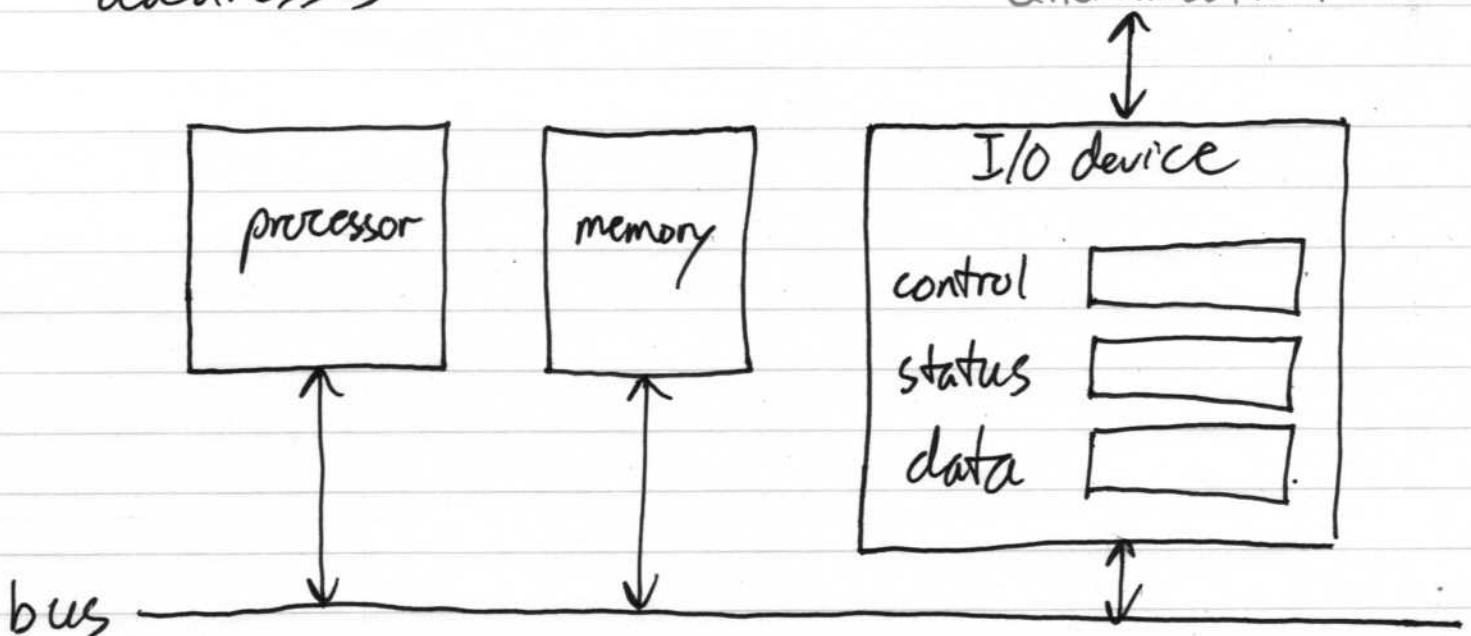
①

1) Memory-Mapped I/O

- the addresses on the bus are shared between memory and peripheral devices
- examples from the LPC1768 memory map:

Address Range	Device
0x0000 0000 - 0x0007 ffff	Flash Memory
0x1000 0000 - 0x1000 7fff	SRAM Memory
0x1fff 0000 - 0x1fff ffff	ROM Memory
0x2009 c000 - 0x2009 8fff	GPIO (general-purpose I/O)
0x4000 c000 - 0x4000 c030	UART 0 (serial port)
0x5000 0000 - 0x5000 0fff	Ethernet

- peripheral devices have registers mapped to specific addresses



e.g. UART - universal asynchronous receiver transmitter ^{serial port}

- data register: Receive Buffer Register (RBR) 0x4000 c000
- control register: Interrupt Enable Register (IER) 0x4000 c004
- status register: Line Status Register (LSR) 0x4000 c014

e.g. code to control UART & interrupts

in "uart.h"


```
#define IER_RBR 0x01 // 0000 00012
#define IER_THRE 0x02 // 0000 00102
#define IER_RLS 0x04 // 0000 01002
```

- these are "bit masks" that give names to bits in the Interrupt Enable Register

in "uart.c"

bit-wise OR
↓

before IER



set bits
0 and 2

```
LPC_UART->IER |= IER_RBR; 00000001
LPC_UART->IER |= IER_RLS; 00000101
```

clears bit
0

```
LPC_UART->IER &= ~IER_RBR; 00000100
```

- setting or clearing these bits controls when the UART interrupts the processor

(3)

- how does LPC_UART → IER address(reference) the IER register at 0x4000c004
- this struct is defined in <LPC17xx.h>:

a union overlays its members in the same location

```
typedef struct {  
    union {  
        __IO uint8_t RBR;  
        __IO uint8_t THR;  
        __IO uint8_t DLL;  
        uint32_t RESERVED0;  
    };  
    union {  
        __IO uint8_t DLM;  
        __IO uint32_t IER;  
    };  
    ...  
    __IO uint8_t LSR;  
} LPC_UART_TypeDef;
```

not a keyword

struct offset

0	RBR/THR/DLL
1	RESERVED0 ↓
2	
3	
4	DLM/↑
5	IER ↓
6	
7	
...	
20	LSR
...	

0x14 → 20

- pointer LPC_UART_TypeDef *LPC_UART is assigned to address 0x4000c000

∴ LPC_UART → IER references 0x4000c004

to do this manually:

LPC_UART = (LPC_UART_TypeDef*) 0x4000c000;

→ Cortex-M3
-in <core_cm3.h>

(input) #define __I volatile const

↑
type modifier that tells the compiler not to optimize out reads or writes to this variable

means that the program cannot assign to this variable (it does not mean that it is unchanging)

(output) #define __O volatile

#define __IO volatile

! the volatile keyword is needed whenever accessing memory-mapped peripherals

2) Polling I/O

- the program repeatedly checks an I/O device's status register until it sees that the device is ready to send or receive data

- in "uart.c"

```
uint8_t UARTReceiveChar(...) {
    ...
    while (!(LPC_UART->LSR & 0x01));
    return (LPC_UART->RBR);
}
```

Annotations:

- `uint8_t`: data type
- `UARTReceiveChar(...)`: function name
- `...`: ellipsis
- `while (!(LPC_UART->LSR & 0x01));`: loop body (labeled "loop body" with an arrow)
- `LPC_UART->LSR`: Line Status Register (labeled "read Line Status Register" with an arrow)
- `volatile`: keyword (labeled "volatile" with an arrow pointing to the register access)
- `0x01`: bit mask (labeled "7 6 5 4 3 2 1 0" with a diagram below)
- `return (LPC_UART->RBR);`: read from data register (labeled "read from data register" with an arrow)
- `RBR`: Receive Buffer Register (labeled "while be 1 if data is received" with an arrow)

Diagram of the LSR register:

7	6	5	4	3	2	1	0
x	x	x	x	x	x	x	?

- this code repeatedly reads the UART's status register until bit 0 is a '1'; then it reads a byte of data and returns it

6

- polling can waste processor time and block other useful execution

e.g. the above polling loop in ARM assembly:

1 cycle	LOOP	LDRB	R1, [R0]	← address of LSR
1 cycle		TST	R1, #0x01	
2 cycles		BEQ	LOOP	
<hr/>				
4 cycles				

- the Cortex-M3 clock rate = 96 MHz
- the "polling rate" = $96 \times 10^6 \frac{\text{cycles}}{\text{s}} / 4 \frac{\text{cycles}}{\text{loop}} = 24 \times 10^6 \frac{\text{loops}}{\text{s}}$

- let's assume that the UART is receiving keyboard data

- a fast typist $\approx 70 \frac{\text{words}}{\text{min.}}$

- average English word length is 5.1 $\frac{\text{char}}{\text{word}}$

- the "data rate" $\approx ?$
 $= \frac{70 \text{ word/min}}{60 \text{ s/min}} \times 5.1 \frac{\text{char}}{\text{word}} = 5.95 \frac{\text{char}}{\text{s}}$

- how many times is the status reg. polled on average per character received?


$$\frac{\text{polls}}{\text{char}} = \frac{\text{polling rate}}{\text{data rate}} = \frac{24 \times 10^6 \frac{\text{loops}}{\text{s}}}{5.95 \frac{\text{char}}{\text{s}}} = 4.03 \times 10^6 \frac{\text{loops}}{\text{char}}$$

- a simple strategy to get more work done:

```
while (true) {  
    do work...  
    check UART status  
}
```


3) Interrupt-Driven I/O

- it may be better to let the peripheral device signal the processor when it needs service (e.g. when the UART receives data)




- three exception types

- ① Interrupt - request for service from a peripheral device
- ② Fault - indicates a problem executing an instruction e.g. divide by zero, bad memory address
- ③ Trap - special instructions that request OS services e.g. SVC in ARM

(8)

- in all cases an exception causes a change in control flow similar to a subroutine call except that the hardware ~~is~~ invokes the subroutine, not the software



- the application code is unaware that it was interrupted
- terminology:
 - in general, exceptions cause "exception handlers" to execute
 - "ISR" is another name for exception handlers for Interrupts
 - but ARM calls ISRs "IRQ Handler"

3.1) Exception Handling Mechanism

[8.2.3, 8.6]

- every exception source, including IRQs, is assigned a "vector" (a number)
- the Vector Table lists the ISR addresses for each vector
- from "startup_LPC17xx.s":

	vector	address of
	0	(holds initial stack pointer)
faults and trap instr. (internal)	1	Reset_Handler runs on startup or reset
	2	NMI_Handler non-maskable interrupt
	...	
	11	SVC_Handler service calls to OS
interrupts (external)	...	
	16	WDT_IRQHandler watchdog timer
	17	Timer0_IRQHandler
	...	
	21	UART0_IRQHandler
	...	

- watchdog timers are used in autonomous systems
 - the application must periodically reset the timer ("kick the dog")
 - if it doesn't and the timer reaches the terminal count it resets the system
 - it is used to recover from system failure

- when an IRQ is received, the processor hardware:

- takes 12 cycles on context switch
- ① looks up the handler address in the vector table
 - ② saves the current execution context on the stack (stack pointer, general-purpose registers, program counter, status register)
 - ③ invokes the exception handler (ISR)

- when the ISR exits, the hardware:


- 10 cycles
- ① restores the execution context from the stack
 - ② resumes the application code

3.2) Multiple Interrupts

[8.3]

- each exception source (IRQ) is assigned a priority

- IRQs of higher priority can interrupt ISRs of lower-priority IRQs ("nested interrupts")



- ISRs of lower-priority IRQs wait for the ISRs of higher-priority IRQs to finish and then run before the application code is resumed

3.3) Interrupt Service Routines (ISRs)

- normally, dealing with the condition needing service clears the interrupt request (IRQ)
 e.g. reading received data from the UART
 e.g. resetting the timer
- if the IRQ is not cleared, the ISR will be re-invoked when it exists
 - sometimes it is necessary to disable IRQs on a peripheral device (by writing to its control register)

! rules for ISRs

- ① keep it short
- ② use in-memory buffers to pass data to/from the application
- ③ if an ISR calls a function, the function must be "re-entrant"

- re-entrant functions can be executed while another instance is executing i.e. can be called concurrently

- they cannot use global or static variables

e.g. ^{not re-entrant} `uint32_t countup() {`

static `uint32_t count;`

`count = count + 1;`

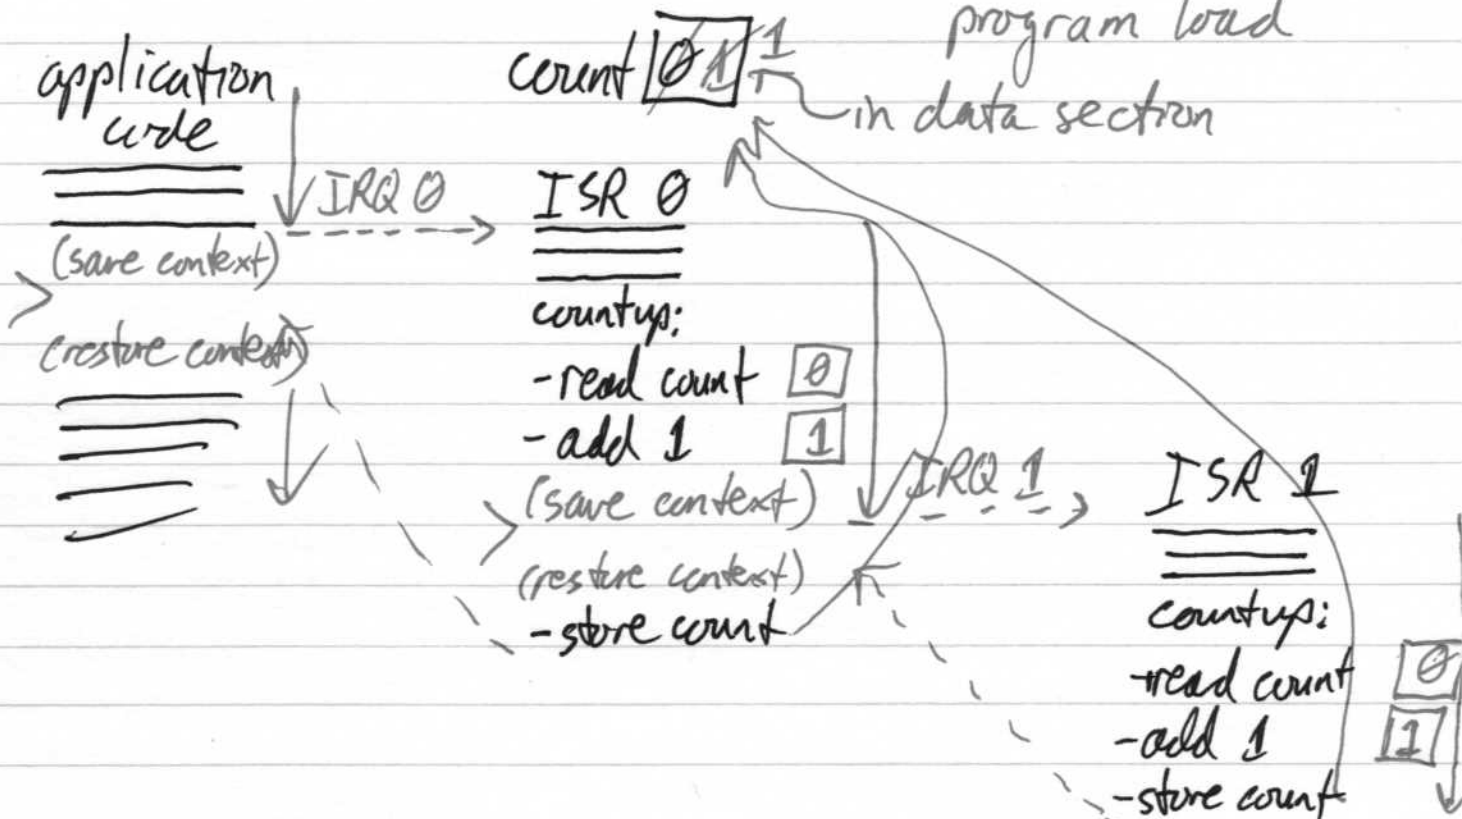
`return count;`

`}`

only one copy that is shared by all invocations of the function

- like a global variable but with local scope

- initialized to 0 on program load



- to use a non-reentrant function in an ISR, you need to disable interrupts

e.g. `--disable_irq();` ← use sparingly

```
    countup();  
} --enable_irq();
```

`--disable_irq()` and `--enable_irq()` are part of the Cortex Microcontroller Software Interface Standard (CMSIS)

- `printf()` and `malloc()` are not re-entrant but as of C11 standard they are "thread-safe" (by adding mutex locks)