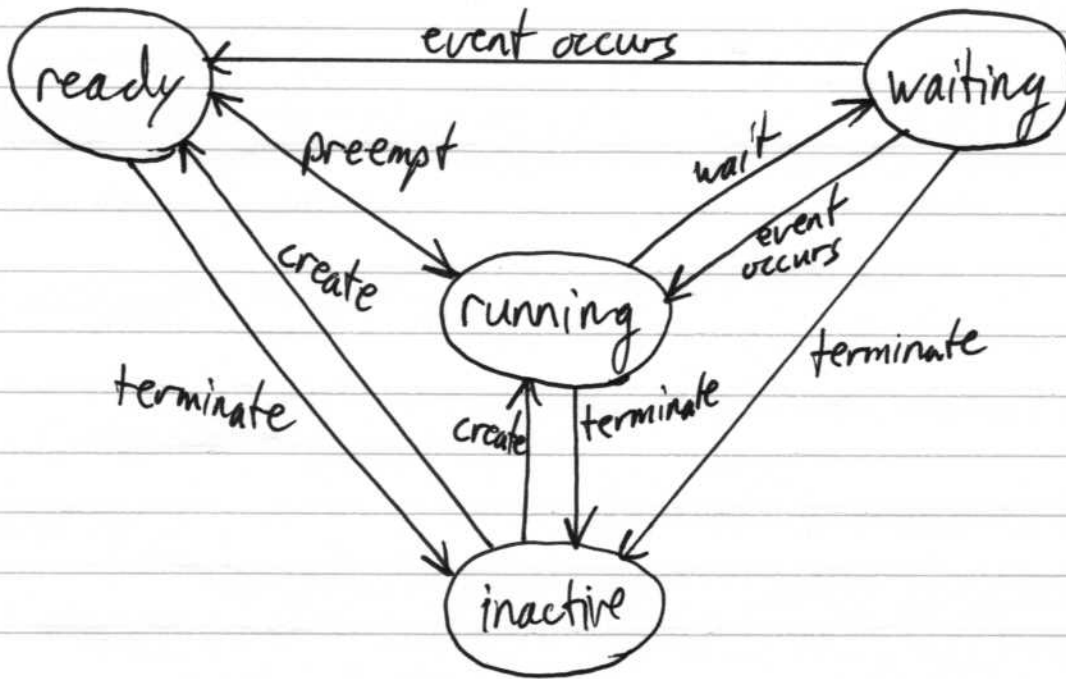


4) Scheduling

[7.3]

4.1) Task States




- the scheduler runs when a task:
 - is created
 - terminates
 - yields e.g. `os Thread Yield()` - tell OS to run another task
 - blocks e.g. `os SignalWait()`, `os Delay()`
 - unblocks e.g. signal received, delay ends
 - finishes its timeslice (default RTX timeslice = 5ms \Rightarrow 200Hz)

4.2) Task Types

[7.2.4.1]

- ① periodic - repeats regularly
e.g. polling I/O
kicking the watchdog timer
- ② aperiodic - can occur any time
- prohibits schedule analysis since it occur any time
and cause indeterminate delays
- ③ sporadic - irregular release(execution) but with
maximum frequency
- permits schedule analysis
e.g. ^{axle} rotation sensor

- tasks can be characterized by:
- T period (or for sporadic task minimum inter-arrival time)
 - C worst-case execution time (WCET)
 - D deadline (relative to release time)



- we will assume $D_i = T_i$ to simplify schedule analysis

- WCET analysis methods

① measurement (dynamic)

- most common method used in industry
- run the program many times with different inputs
- runtime is measured with a profiling tool such as gprof or hardware timers
- this can underestimate WCET so safety margins are added

② static analysis


- flow-control analysis

e.g.

```
do {
  if (x > 0) {
    ...
  } else {
    ...
  }
} while (i != j);
```

- estimate loop bounds and longest path

- use low-level analysis to estimate time of basic blocks or individual instructions
- flow-control and timing estimates are combined to calculate the WCET
- prone to over-estimating WCET



③ hybrid of static analysis and measurement of basic blocks

4.3) How to Evaluate a Scheduler

- ① minimize missed deadlines or minimize lateness
- ② maximize processor utilization

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq 1 \quad \text{percentage of time that the processor is busy}$$

- ③ minimize scheduler overhead (time to decide next task to execute)

4.4) Non-preemptive Scheduling

[7.3]

- advantage: minimizes scheduler overhead

4.4.1) Timeline Scheduling

[7.3.1]

a.k.a. superloop

- tasks execute in a fixed order which repeats indefinitely
- the schedule is created offline

e.g. task set

	<u>T (period)</u>	<u>C (exec)</u>
τ_1	18	8
τ_2	30	10
τ_3	45	10

 $(D_i = T_i)$

- check processor utilization

$$U = \frac{8}{18} + \frac{10}{30} + \frac{10}{45} = 1 \leq 1 \quad \checkmark$$

(16)

- schedule length = LCM (least common multiple) of task periods

- LCM calculation techniques

$$\textcircled{1} \quad \text{LCM}(a, b) = \frac{a * b}{\text{GCD}(a, b)}$$

GCD = greatest common divisor

$$\text{LCM}(a, b, c) = \text{LCM}(\text{LCM}(a, b), c)$$

$\textcircled{2}$ decompose into prime factors

- find smallest prime divisor
- divide period by divisor and repeat

eg.	$18 \div 2 = 9$	$30 \div 2 = 15$	$45 \div 3 = 15$
	$9 \div 3 = 3$	$15 \div 3 = 5$	$15 \div 3 = 5$
	$3 \div 3 = 1$	$5 \div 5 = 1$	$5 \div 5 = 1$
	$18 = \textcircled{2} \times \textcircled{3^2}$	$30 = 2 \times 3 \times \textcircled{5}$	$45 = 3^2 \times 5$


then multiply the prime factors of highest power

eg. $\text{LCM} = 2 \times 3^2 \times 5 = \underline{90}$

- create schedule using Earliest Deadline First (see handout)


	T_i	C_i
τ_1	18	8
τ_2	30	10
τ_3	45	10

Method: non-preemptive EDF



	T_i	C_i
τ_1		
τ_2		
τ_3		

Method: _____



(17)

- implementation: store the order in an array and the scheduler moves to the next element (task) each time
- scheduler overhead: $O(1)$

$$\text{array length} = \sum_{i=1}^n \frac{\text{LCM}}{T_i} = \frac{90}{18} + \frac{90}{30} + \frac{90}{45} = 10$$

- disadvantage: the storage required for the schedule can be ^{large}

e.g. 5 tasks with deadlines $T_1=18, T_2=30, T_3=45, T_4=22, T_5=23$

$$\text{schedule length} = 2 \times 3^2 \times 5 \times 11 \times 23 = 22770$$

$$\text{array length} = \frac{22770}{18} + \frac{22770}{30} + \frac{22770}{45} + \frac{22770}{22} + \frac{22770}{23} = 4555$$

- to shorten schedule and array length, you can reduce some task periods, as long as it is still schedulable ($U \leq 1$)

$$\text{e.g. } T_1=18, T_2=30, T_3=45, T_4=\underline{20}, T_5=20$$

$$L = 2^2 \times 5$$

$$\text{schedule length} = 2^2 \times 3^2 \times 5 = 180$$

$$\text{array length} = 180 \left(\frac{1}{18} + \frac{1}{30} + \frac{1}{45} + \frac{1}{20} + \frac{1}{20} \right) = 38$$

4.5) Preemptive Schedulers


[7.4]

- tasks are suspended in mid-execution to allow another task to run


4.5.1) Round-Robin Preemptive

round robin

- all tasks have equal priority
- tasks switch ~~on the SysTick interrupt~~ ^{at the end of their timeslice} each



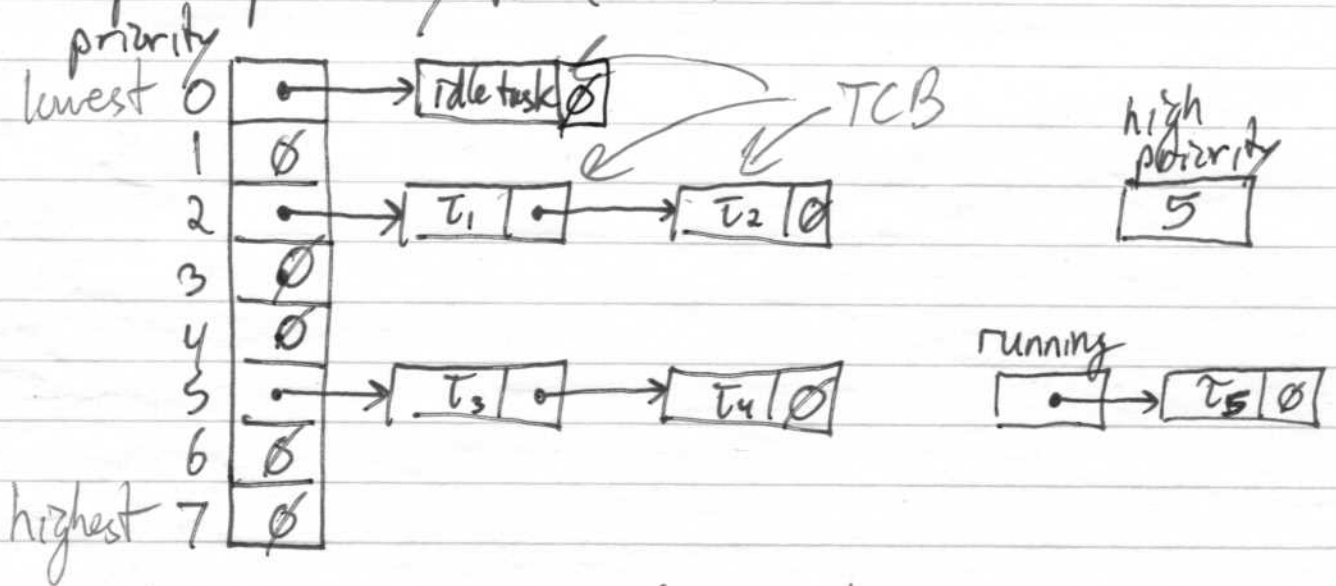
- implementation: array or a linked-list queue
 - enqueue current task at the tail and dequeue next task to run from the head
 - we can use the link pointers in the task control blocks (TCBs)



- scheduling overhead is $O(1)$
- disadvantages: no priorities and no deadlines

4.5.2) Fixed-Priority Preemptive Scheduling [7.4.5]

- task priorities are static (don't change)
- run the ready tasks with highest priority in round-robin fashion
- implementation: an array of linked-list queues, one per priority level




- when a timeslice (sys tick) ends, enqueue the current task at the back of the queue and deque next front the front
- adding a task is $O(1)$
- what happens when the high-priority queue empties?

① linear scan for next non-empty queue
= $O(n)$

or ② use a bit vector to represent non-empty queues

e.g.



- the first set bit can be found with an instruction such as ARM CLZ (count leading zeroes) instruction
- can also be done in software in constant time using deBruijn sequences

- advantage: $O(1)$ scheduling overhead \Rightarrow widely used

- disadvantages:
 - no deadlines
 - frequent scheduling / context switches

- Rate Monotonic scheduling (RM) [7.4.5.2]

- a way to assign ^{fixed} priorities based on deadlines
- priority \propto execution rate of task
- (shorter period \Rightarrow higher priority)
- (higher frequency)

- schedulability test:

a set of n tasks are schedulable by RM if

$$\prod_{i=1}^n \left(1 + \frac{C_i}{T_i}\right) \leq 2$$

- the test to accept a new task

- this can be performed in $O(1)$ (advantage)

4.5.3) Rate Monotonic Scheduling [7.4.5.2]

- it is a way to assign ~~prop~~ priorities for F.P.P. based on deadlines
 - priority is relative to the task frequency (shorter period \Rightarrow higher priority)

- schedulability test:
a set of n tasks are schedulable by RM if

$$\prod_{i=1}^n \left(1 + \frac{C_i}{T_i}\right) \leq 2$$

- the test to accept a new task can be performed in $O(1)$ time

e.g.

	T_i	C_i	$(D_i = T_i)$
τ_1	10	2	
τ_2	10	3	
τ_3	6	1	

- check processor utilization

$$U = \frac{2}{10} + \frac{3}{10} + \frac{1}{6} = 0.67 \leq 1 \quad \checkmark$$


- RM schedulability

$$\prod_{i=1}^n \left(1 + \frac{C_i}{T_i}\right) = \left(1 + \frac{2}{10}\right) \left(1 + \frac{3}{10}\right) \left(1 + \frac{1}{6}\right) = 1.82 \leq 2 \quad \checkmark$$

Method: RM with FPP


	T_i	C_i
τ_1	20	2
τ_2	10	3
τ_3	6	1

} tied for lowest priority.
 - assume time slice = 1 time unit
 - highest priority



	T_i	C_i
τ_1		
τ_2		
τ_3		

Method: _____



4.5.4) Preemptive EDF

[7.4.4.1]

- priorities are dynamic (deadline is the priority)
- the scheduler runs the task with the closest deadline
- if a task is released with an earlier deadline it preempts the current task
- no timeslices - no SysTick interrupt
- implementation: two priority queues
 - ① "ready" queue - all ready tasks, sorted by deadline
 - ② "waiting" queue - waiting tasks sorted by next release
- the scheduler overhead is $O(\log n)$ assuming that minheaps are used
- advantages:
 - optimal scheduling algorithm
 - has the fewest scheduler invocations and fewest context switches
- disadvantages:
 - has a complex schedulability test (which is needed to decide if a new task can be admitted)