

Synchronization

①

1) Semaphores


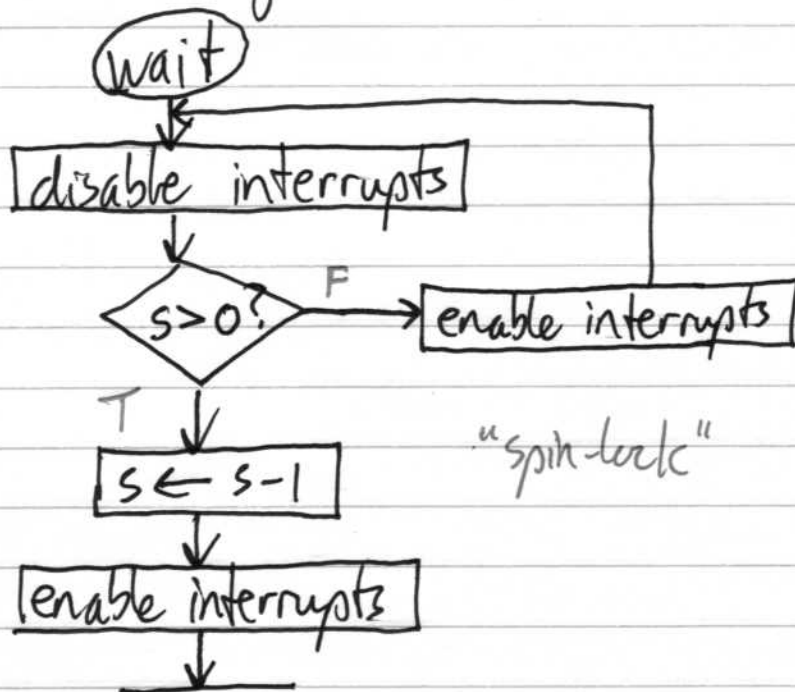
[9.5]

- a semaphore is a counter with 3 functions

① init - initialize counter value S

② wait - tries to decrement counter; wait if $S=0$

③ signal - increments counter



- read-modify-write sequences are protected by disabling interrupts
- the spin lock must enable interrupts each iteration to allow the SysTick interrupt so the scheduler can run (so other tasks can run)

(2)

6.1) Example Semaphore Implementation without RTX

```
typedef uint32_t sem_t;
```

```
void init(sem_t *s, uint32_t count) {
    *s = count;
}
```

```
void wait(sem_t *s) {
```

```
    __disable_irq();
```

← CMSIS

```
    while (!(*s > 0)) {
```

```
        __enable_irq();
```

```
        __disable_irq();
```

← this gives scheduler a chance to run

```
    }
    (*s)--;
```

```
    __enable_irq();
}
```

```
void signal(sem_t *s) {
```

```
    __disable_irq();
```

```
    (*s)++;
```

```
    __enable_irq();
}
```

2) Mutual Exclusion

- protect code that accesses shared data (critical code) by using a semaphore as a lock

- steps:

- ① initialize the semaphore lock to 1 (open)
- ② invoke `wait()` to acquire the lock ~~when~~ before entering the critical section
- ③ invoke `signal()` to release the lock when leaving the critical section

e.g. `sem_t lock;`

```
int main(void) {
    init(&lock, 1);
    ... start task 1 and task 2
}
```

	<code>void task1(void) {</code>	<code>void task2(void) {</code>
acquire	<code>wait(&lock);</code>	<code>wait(&lock);</code>
	- execute in critical section	- execute in critical section
release	<code>signal(&lock);</code>	<code>signal(&lock);</code>
	<code>...</code>	<code>...</code>
	<code>}</code>	<code>}</code>

- only the lock "owner" should release it

3) Condition Variable

- one task signals another task that an event has occurred; then the other task can proceed

- steps:

- ① initialize the condition variable semaphore to 0
- ② one task invokes wait() to wait for the event
- ③ the other task (or ISR) invokes signal() to indicate that the event occurred

e.g. sem-t cond; condition variable

```
int main(void) {
    init(&cond, 0);
    - start task1 and task2
}
```

```
void task1(void) {
    // wait for event
    wait(&cond);
    - continue execution
    ...
}
```

```
void task2(void) {
    // detect event
    signal(&cond);
    - continue execution
    ...
}
```

4) Task Rendezvous

- synchronize two tasks to perform work at the same time

- steps:

- ① initialize two semaphores to 0
- ② each task signals the other task
- ③ both tasks wait for the other's signal

e.g. sem_t s1, s2;

```
int main(void) {
    init(&s1, 0);
    init(&s2, 0);
    - start task1 and task2
}
```

```
void task1(void) {
```

```
    ...
    // rendezvous
    signal(&s2);
    wait(&s1);
    ...
```

```
}
```

```
void task2(void) {
```

```
    ...
    // rendezvous
    signal(&s1);
    wait(&s2);
    ...
```

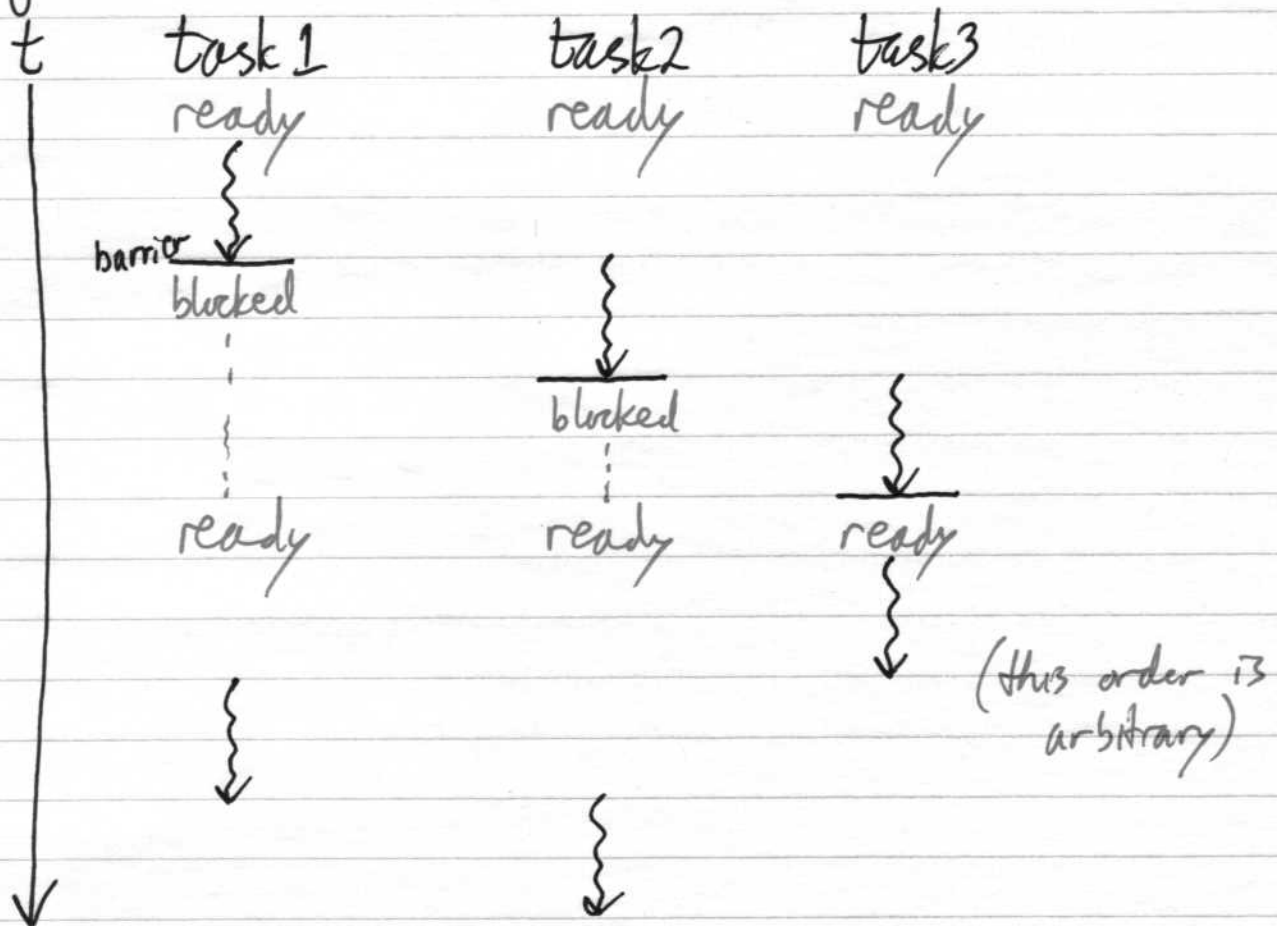
```
}
```

6

5) Barrier Synchronization

- all tasks must arrive at the barrier before any proceed past it
i.e. it is an n -task rendezvous

eg. 3 task barrier



- we can try to solve with n semaphores

eg. task1 s1 task2 s2 task3 s3

 signal(&s2); signal(&s3); signal(&s1);

 wait(&s1); wait(&s2); wait(&s3);

okay?
NO

- see further A3

⑧

- general solution [The Little Book of Semaphores, Downey]

- 1 barrier semaphore
- 1 mutex semaphore
- 1 count variable

- steps:

① initialize barrier = 0, mutex = 1, count = 0

make tasks wait

track arrivals

protect accesses to count

② each task:

- acquire mutex
- increment count
- release mutex

③ last task signals the barrier

④ each task:

- wait on barrier
- signal barrier

} "turnstile"

- implementation {

```
sem_t barrier, mutex;
uint32_t count;
```

```
void barrierInit(void) {
    init(&barrier, 0);
    init(&mutex, 1);
    count = 0;
}
```

```
(critical section) void barrierSync(void) {
    wait(&mutex);
    count++;
    signal(&mutex);
    if (count == n)
        signal(&barrier);
    wait(&barrier);
    signal(&barrier);
}
```

```
int main(void) {
    barrierInit();
    // start n tasks
}
```

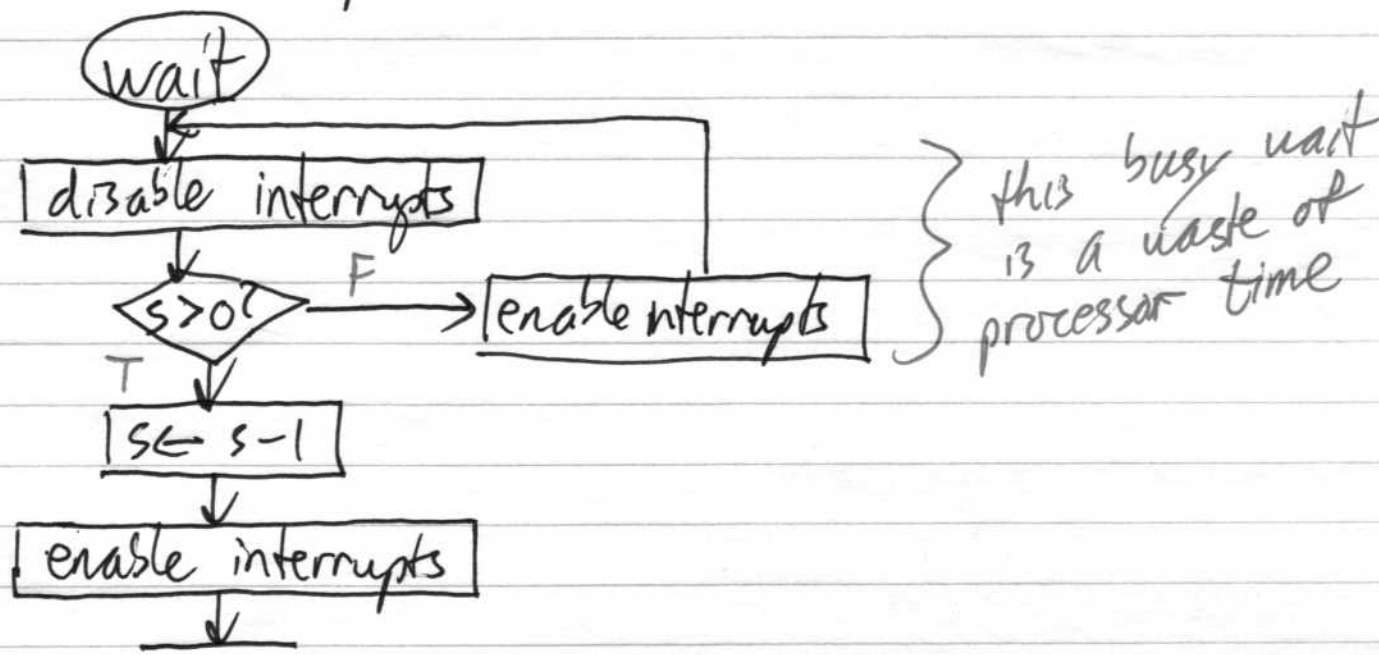
```
void task(void) {
    ...
    barrierSync();
    ...
}
```

- value of the barrier semaphore will be in $[1, n]$ after all have synced

- is this barrier re-usable?

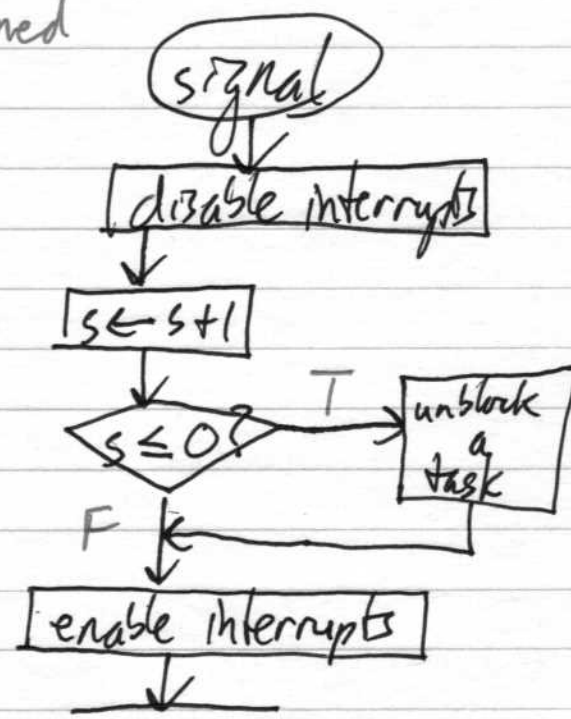
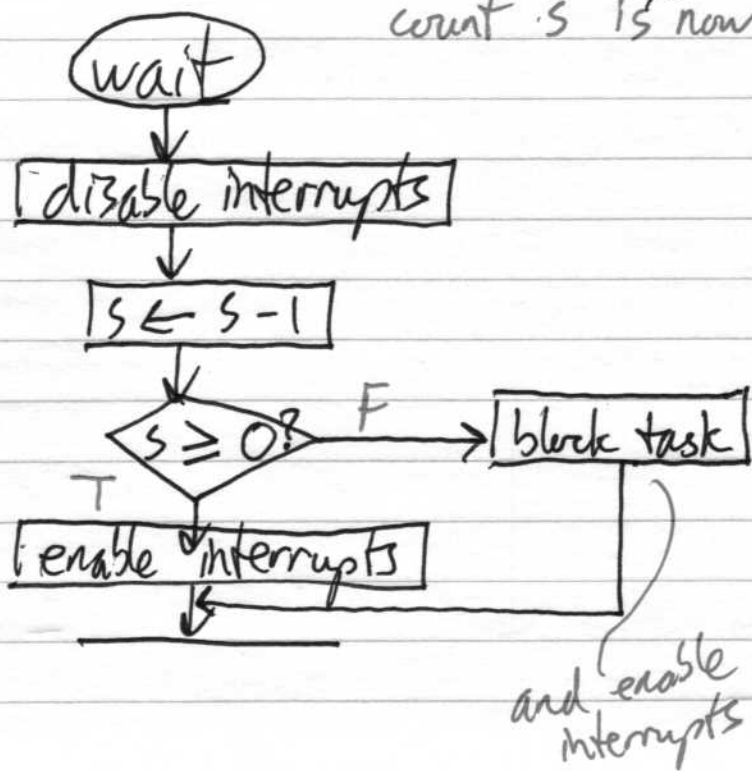
6) Example Implementation of a Blocking Semaphore

- recall the "spin lock" in wait()



- a better implementation is to block the waiting task (allowing another task to run), and signal() selects one block task (if any) to unblock



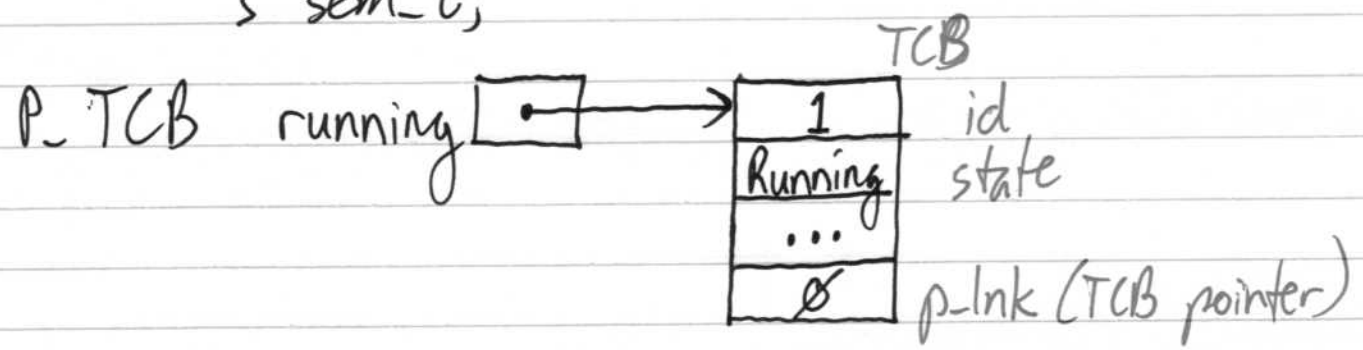
count S is now signed



- each semaphore now has a wait list associated with it

e.g.

```
typedef struct {
    int32_t s;
    P_TCB waitList; task control block pointer
} sem_t;
```



- suppose that Task 1 waits on sem1 which is not available ($s \leq 0$)
- Task 1 is moved to the tail of sem1.waitList
- Task 1's state changes to Waiting
- Task 2 is moved from the head of the ready list to the running "list"
- Task 2's state changes to Running

7) CMSIS-RTX Semaphores

- two types: general semaphore (counting semaphore)
mutual exclusion semaphore (mutex)
- general semaphore use

```
#include <cmsis_os.h>
```

```
osSemaphoreId sem;
osSemaphoreDef(sem);
sem = osSemaphoreCreate(osSemaphore(sem), 0);
    - returns NULL on error
osSemaphoreRelease(sem);
    - returns osStatus (osOK or error code) "signal"
osSemaphoreWait(sem, osWaitForever);
    - returns # of tokens (0 = no token, -1 = error)
    - cannot be called by exception handlers (ISRs)
```

initial count

timeout in ms (0 = no wait)

- mutex use

- semaphore count is 1 or 0

```
osMutexId lock;
osMutexDef(lock);
lock = osMutexCreate(osMutex(lock)); initial count = 1
```

- returns NULL on error

neither can be called by an exception handler

```
{ osMutexWait(lock, osWaitForever); "acquire"
  osMutexRelease(lock); }
```

- RTX mutexes are recursive

- the same task can acquire the mutex multiple times without blocking
- it must release the mutex an equal number of times
- it is a convenience feature (not necessary)
- the problem is that it can encourage holding the mutex too long
- you should hold a mutex for as little time as possible to maximize concurrency
c.f. David Butenhot re: recursive mutex

- RTX mutexes check if the task owns the mutex (i.e. has acquired it) on release

- CMSIS-RTOS v1 - no priority inheritance
- CMSIS-RTOS v2 - priority inheritance

8) Priority Inversion

- priority inversion occurs when a higher priority task is indirectly blocked from running by a lower priority task
- it requires 3+ tasks to occur

e.g. Mars Pathfinder Reset Problem

[10.5.1]

- used VxWorks RTOS

- 3 tasks:

① bus .