

3) Double Buffering

- 2 message buffers allow parallel production and consumption
- used in graphics to avoid "tearing"
- we need a pair of semaphores (empty, full) per buffer

e.g.
buffer 0

		buffer 0		buffer 1		sum	
		empty	full	empty	full	empty	full
<u>producer</u>	<u>consumer</u>	1	0	1	0	2	0
wait b0 empty	wait b0 full	0	0	1	0	1	0
produce	block						
signal b0 full		0	1	1	0	1	1
wait b1 empty	unblock	0	0	0	0	0	0
produce	consume						
signal b1 full	signal b0 empty	1	0	0	1	1	1
wait b0 empty	wait b1 full	0	0	0	0	0	0
produce	consume						
signal b0 full	signal b1 empty	0	1	1	0	1	1
⋮	⋮						


- we can combine the "empty" semaphores and initialize to 2
- we can combine the "full" semaphores and initialize to 0
- as long as the producer(s) and consumer(s) use the buffers in the same order

(4)

- we can statically allocate the buffers and pass pointers to them to the tasks

① statically allocated storage

② array of buffer pointers



e.g. two 512B buffers

```
const uint32_t bufSize = 512;
① uint8_t bufData[bufSize * 2];
② uint8_t *bufPtr[2] = { bufData, bufData + bufSize };
                        &bufData[bufSize]
```

```
sem_t empty, full;
init(&empty, 2);
init(&full, 0);
```

producer

```
uint32_t index = 0;
```

```
while(true) {
```

```
    wait(&empty);
```

```
    ... fill buffer at bufPtr[index]
```

```
    signal(&full);
```


```
    index = (index + 1) % 2;
}
```

consumer

see assignment 4

4) Bounded-Buffer Problem

- generalizes double-buffering to n buffers
- "bounded" refers to the fact that the buffer is not infinite
i.e. producers must wait on empty buffers



- solution: use an array-based queue of buffers (or buffer pointers)
- two semaphores: $\text{empty} = n$
 $\text{full} = 0$
- two indices: $\text{tail} = 0$ for producers
 $\text{head} = 0$ for consumers
- two mutexes: tailMutex protects tail index updates by multiple producers
 headMutex protects head index updates by multiple consumers
- see assignment 4

5) CMSIS-RTX Queues

- two types:

① `osMessageQ` - passes `uint32_t` or `void *` between threads.
- user manages storage

② `osMailQ` - passes pointers to storage managed by the queue


e.g. bounded-buffer solution using `osMessageQ`
- 8 512 B buffers, 2 producers, 2 consumers

```
const uint32_t bufSize = 512;
const uint32_t n = 8;
uint8_t bufData[n * bufSize];
```

} data segment

```
osMessageQDef(emptyQ, n, uint8_t *);
osMessageQId(emptyQId);
osMessageQDef(fullQ, n, uint8_t *);
osMessageQId(fullQId);
```

queue name queue size element type



(7)

```
int main(void) {
```

```
    osKernelInitialize();
```

```
    emptyQId = osMessageCreate(osMessageQ(emptyQ),  
                                NULL); always NULL
```

```
    fullQId = osMessageCreate(osMessageQ(fullQ), NULL);
```

```
    osKernelstartcreate();
```

```
    for (uint32_t i=0; i<n; i++)
```

```
        osMessagePut(emptyQId,  
                    (uint32_t)(bufData + i*bufSize), osWaitForever);
```


\Downarrow
 $\&\text{bufData}[i*\text{bufSize}]$ *args*

```
        osThreadCreate(osThread(producer), NULL);
```

```
        osThreadCreate(osThread(consumer), NULL);
```

```
}
```

(8)



```
void producer (void const *args) {
    while (true) {
```

```
        osEvent event = osMessageGet(emptyQId, osWaitForever);
```

```
        uint8_t *buf =
            event.value.p;
```

```
        ... fill buffer pointed to
            by buf
```

```
typedef struct {
    osStatus status;
    union {
        uint32_t v;
        void *p;
        int32_t signals;
    } value;
} osEvent;
```

```
    osMessagePut(fullQId, (uint32_t)buf, osWaitForever);
}
```

- consumer is the mirror image of the producer


(9)

6) Buffer Sizing

- buffers allow production rates $P(t)$ to temporarily exceed consumption rates $C(t)$ for limited periods of time

- engine oil analogy


[LaPlante, 1997]



- the funnel buffers the oil so none is lost (if it is big enough)

6.1) Constant $P(t)_{\text{burst}}$ and Constant $C(t)$

- if $P > C$ for limited time T , how big should the buffer be?



- data produced during burst = PT
- data consumed during burst = CT
- buffer size $B = PT - CT = (P - C)T$

e.g. I/O device produces data at 9600 B/s for a 1 s burst every 20 s

- a consumer task consumes this data at 800 B/s

$$B = (9600 \text{ B/s} - 800 \text{ B/s}) \times 1 \text{ s} = 8800 \text{ B}$$

\therefore buffer size should be ≥ 8800 bytes


- this assumes that there is time to drain the buffer between bursts

$$t = 8800 \text{ B} / 800 \text{ B/s} = 11 \text{ s} \quad \begin{matrix} < 19 \text{ s} \\ \cancel{19 \text{ s}} \end{matrix} \quad \checkmark$$

7) Queuing Theory

[15.2]

- models the real world as client/server systems
e.g. clients lining up for service at a bank
requests arriving at a web server
- clients arrive and wait in a queue to be served
(producer = client, consumer = server)



$$\rho = \text{load factor} = \frac{\lambda}{\mu}$$

- the arrival and service rates can be:

① deterministic 'D'


- clients arrive with regular interval $T = \frac{1}{\lambda}$
- server time = $\frac{1}{\mu}$

② Markov 'M'

- a Markov process is a memoryless stochastic process

↑
a sequence of random variables

↑
time between each arrival / service is independent of the previous time



probability density function
- probability of the next event occurring at time t

- Kendall notation for queuing models:

arrival model	/	service model	/	n
↓		↓		↓
D or M		D or M		# of servers

7.1) M/M/1 Queues

- equations:

average clients in system $L = \frac{\rho}{1-\rho}$

queued + the
client being served

average time in system $W = \frac{1/\mu}{1-\rho}$

time waiting + time
being served

Little's Rule $L = \lambda W \left(= \lambda \frac{1/\mu}{1-\rho} = \frac{\lambda/\mu}{1-\rho} = \frac{\rho}{1-\rho} \right)$